

Hyperparameter Optimisation of an Adversarial Neural Network in the tW channel at 13 TeV with ATLAS

Christian Kirfel

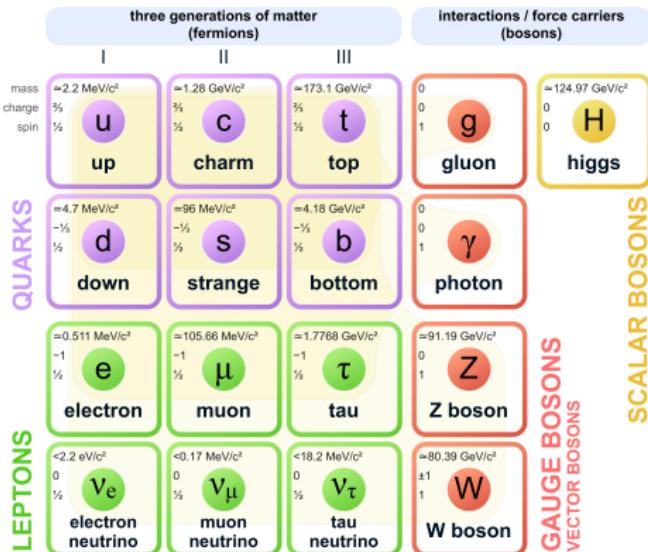
31. März 2019

Outline

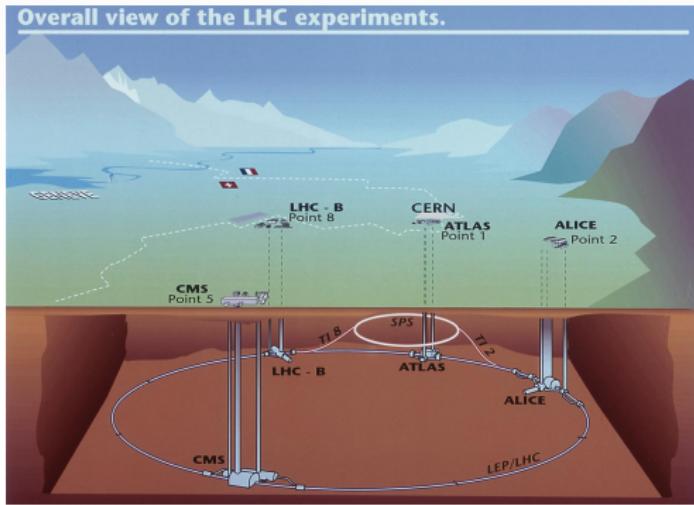
- The Standard Model of particle physics
- The LHC and the ATLAS detector
- Motivating the tW channel
- Neural networks as a possible solution
- Results for different approaches of adversarial neural networks

The standard model of particle physics

Standard Model of Elementary Particles



The Large Hadron Collider - LHC

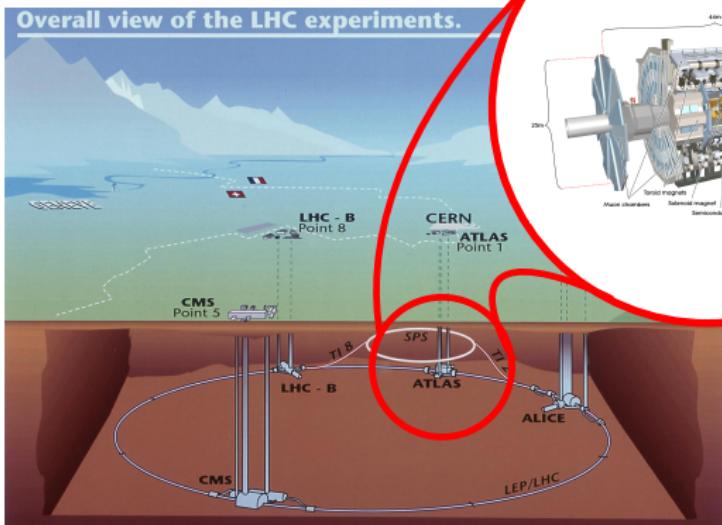


[6]

$$E_{CMS} = 13 \text{ TeV}$$

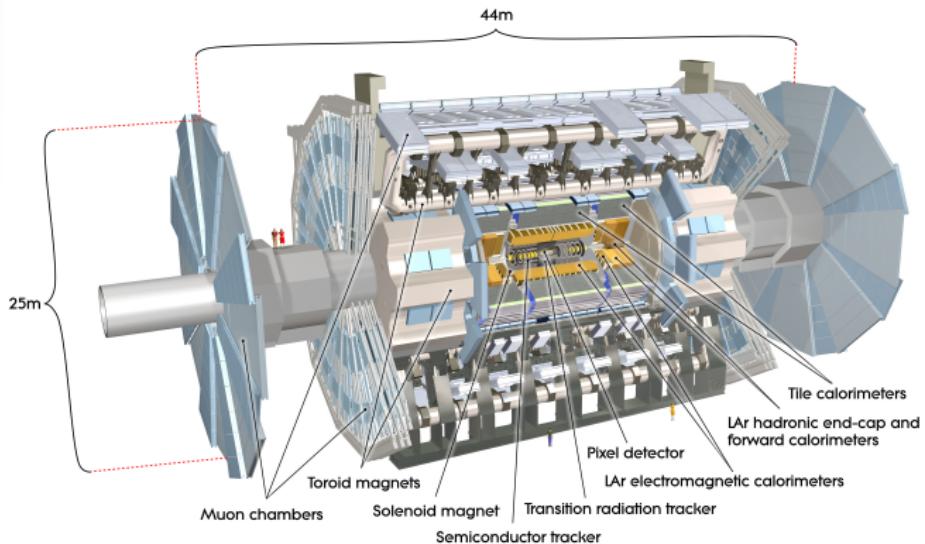
$$\mathcal{L} = 1,5 \cdot 10^{34} \text{ cm}^{-1} \text{ s}^{-2}$$

The ATLAS experiment

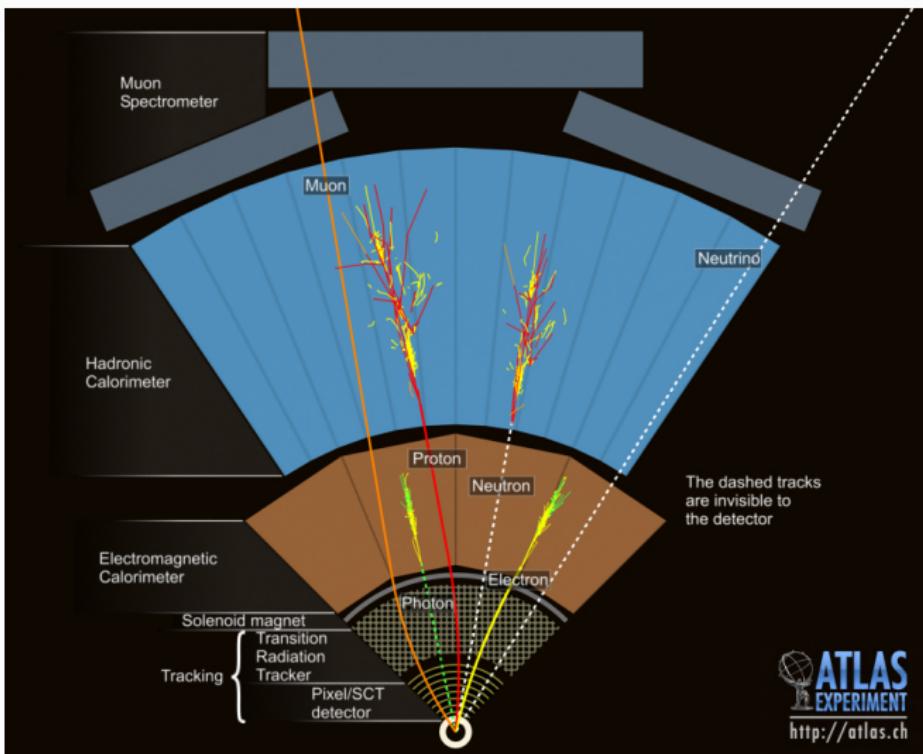


[6]

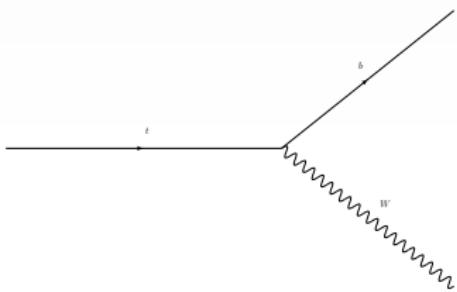
ATLAS - a general purpose detector



ATLAS - object identification



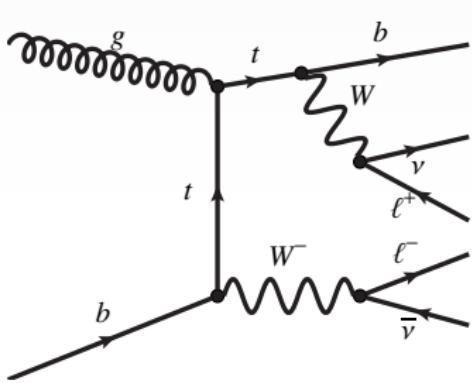
Top physics



- $m \sim 173 \text{ GeV}/c_0^2$
- $\tau \sim 5 \cdot 10^{-25} \text{ s}$

- Decays almost always into a top-quark and a W -boson
- Lifetime shorter than typical time of hadronisation
- No bound states

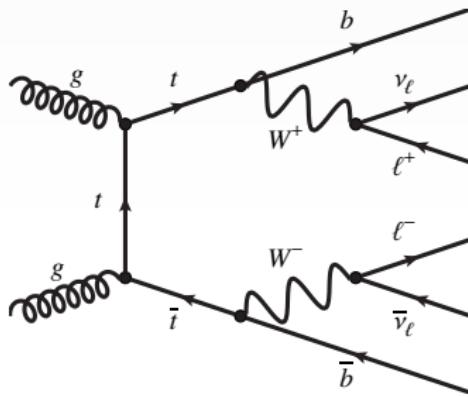
tW to tt} separation at LO



tW decay

$$\sigma_{tW} \sim 71.7 \text{ pb}$$

Final state: 2 W, 1 b

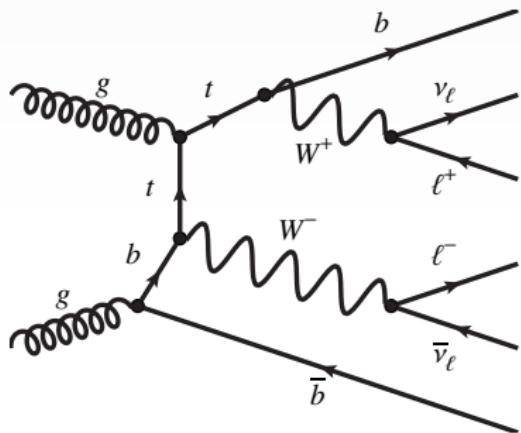


tt} decay

$$\sigma_{t\bar{t}} \sim 832 \text{ pb}$$

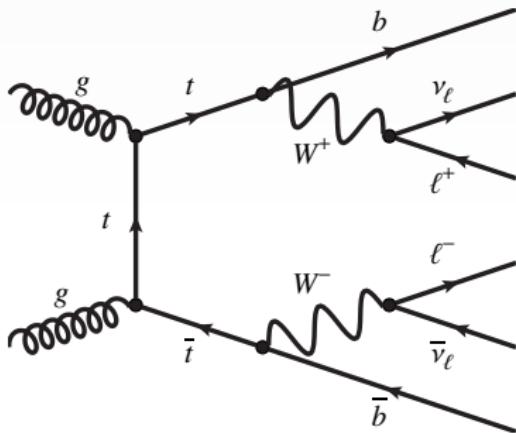
Final state: 2 W, 2 b

tW to $t\bar{t}$ separation at NLO



tW decay at NLO

Final state: 2 W, 2 b



$t\bar{t}$ decay

Final state: 2 W, 2 b

Challenge 1 - Signal to background separation

Problem 1

- Separation of signal to background
- Signal: tW_DR , tW_DS
- Background: $t\bar{t}$

Classic approach

Applying cuts

Alternative

- Machine Learning
- In particular: Classifying neural network

Artificial neural network

Neural Networks - Processing information

Human senses

- Extract relevant information
- Not possible for machines

Human brain

- Web of neuron cells
- A neuron cell takes input from other cells
- Combination of signals creates a result

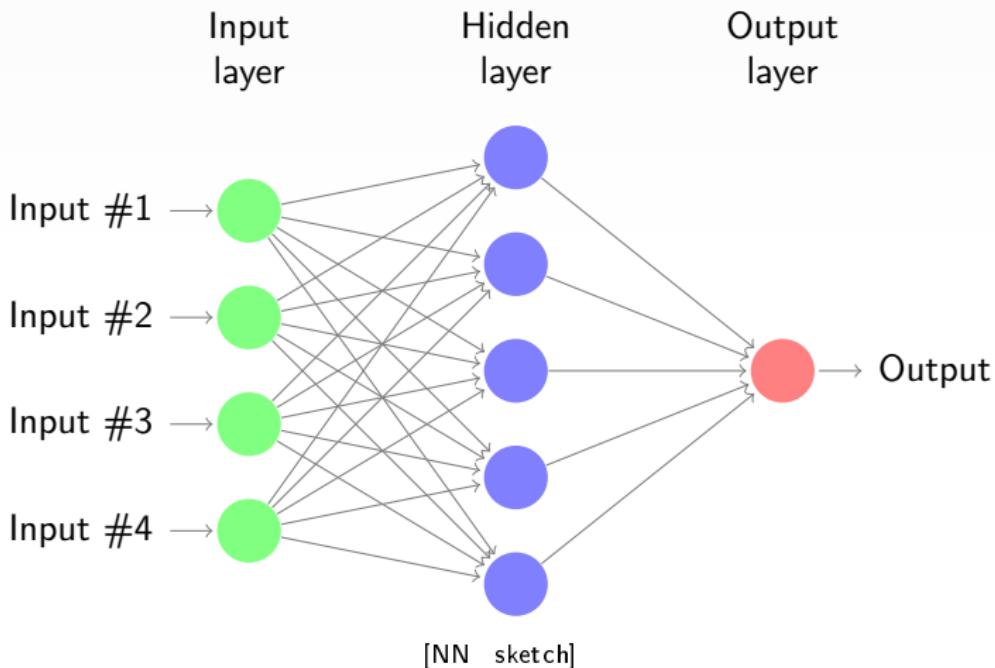
Input variables

- Preprocessed by human user
- e.g. kinematic variables of an event

Net of nodes

- Nodes are simple processors
- Connected by linear function (weight and bias)
- Combines input using an activation function

Neural network structure



Neural Networks - Choosing the next step

Evaluation of an action

- Simple perceptions: pain, satisfaction
- Expected or desired outcome

Deciding for a better action

- Trial and error
- Learning from experience

Loss function

- Supervised learning: compare to the desired outcome
- Loss is an estimator for a model's quality

Optimisation

- Propagate the impact of the parameters on the loss backwards
- Adjust parameters to achieve a model representing a smaller loss

Challenge 2 - Sensitivity to systematic uncertainties

Problem 2

- Minimal sensitivity to the systematic uncertainty
- Nominal: tW_DR , ($t\bar{t}$)
- Systematic: tW_DS

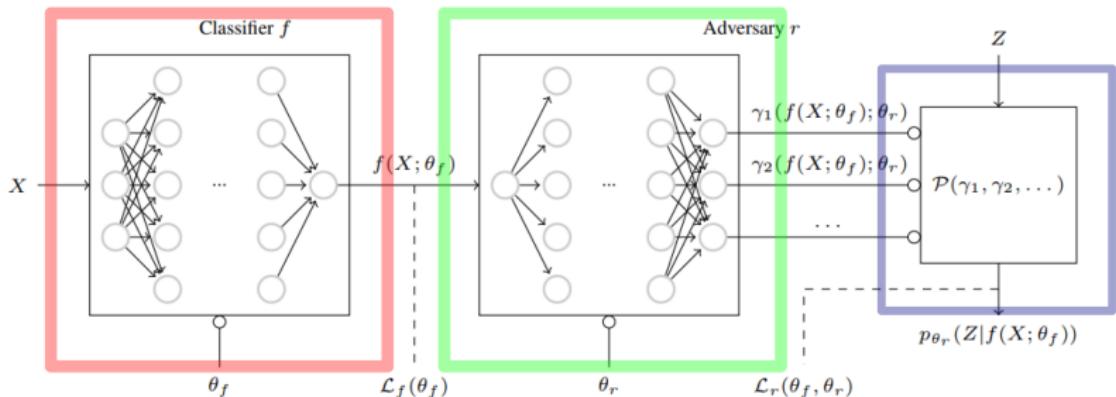
Presented solution

- Addition of a second classifier for nominal to systematics separation
- Bad performance → low sensitivity

Implementation

- Feed output into a second net
- Iterative training
- Combined loss function → Minimax problem

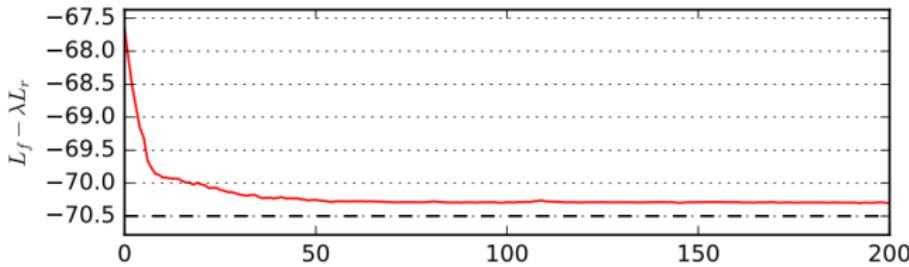
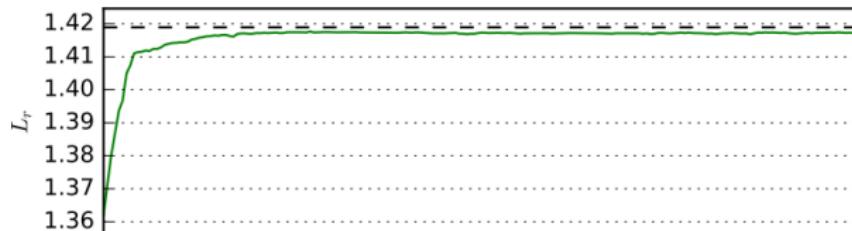
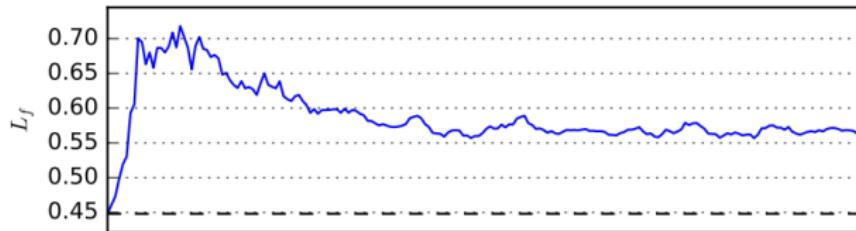
Adversarial neural network



[3]

$$\mathbb{E}(\theta_f, \theta_r) = \mathcal{L}_f(\theta_f) - \lambda \mathcal{L}_r(\theta_f, \theta_r)$$

Expected ANN losses



Classifier training

Setup of the classifier

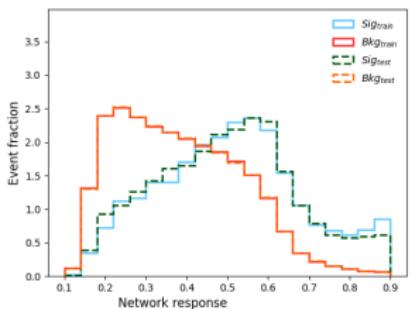
Hyper-parameter scan results

- Input: 14 variables motivated by a BDT variable scan.
- Hidden layers: 6 elu layers \times 128 nodes each
- Output layer: 1 sigmoid node
- Optimisation: SGD, Learning rate = 0.06, momentum = 0.3, no nesterov, no decay
- Duration: 600 epochs

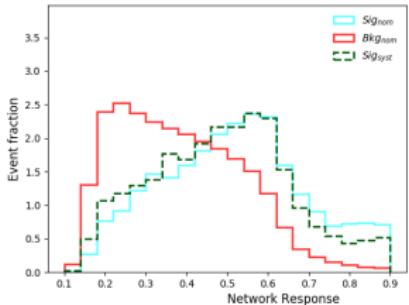
Variables

- Motivated by a BDT significance scan
- Tested against simple kinematic variables

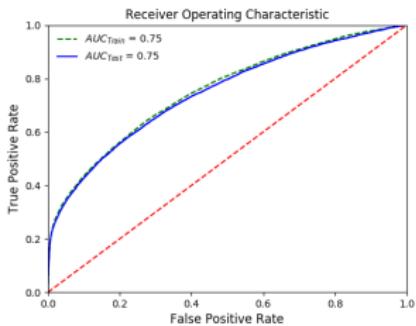
Simple network results



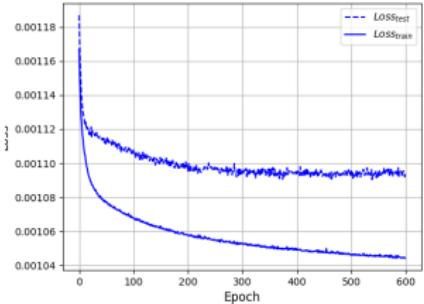
(a) Separation



(b) Systematics



(c) ROC curve



(d) Losses

Adversarial training

Approach 1

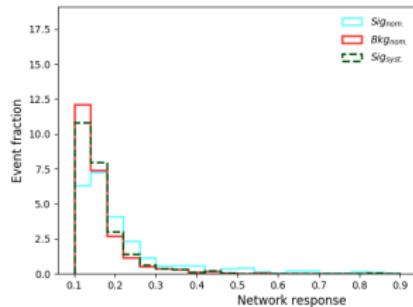
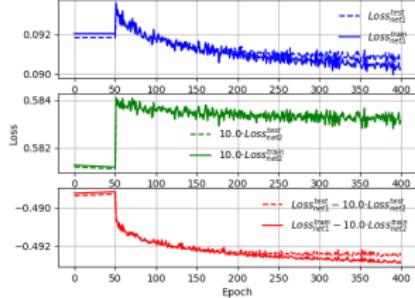
Run 1

- Same hyper-parameters for both networks
- Adversary uses only 3 layers

Setup

- Input: 14 variables motivated by a BDT variable scan.
- Hidden layers: 6(3) elu layers \times 128 nodes each
- Output layer: 1 sigmoid node
- Optimisation: SGD, Learning rate = 0.06, momentum = 0.3, no nesterov, no decay
- Duration: 400 iteration
- $\lambda = 10$

Approach 1 - Run 1 - Results



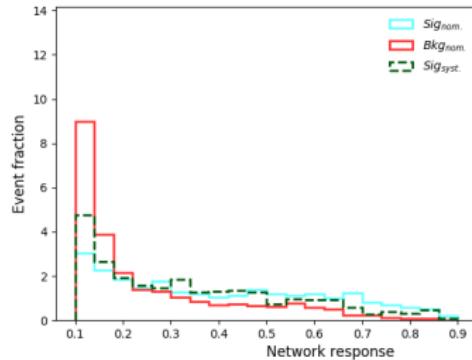
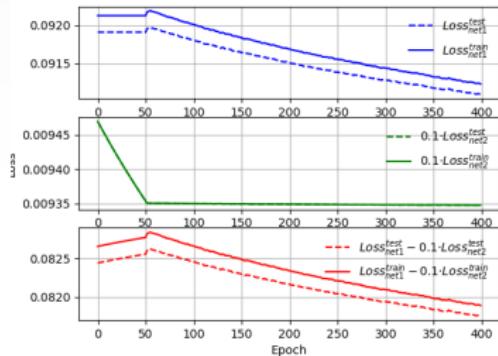
- Both losses decreasing
- Only combined losses depend on λ
- Bad separation
- High sensitivity to systematics

Approach 1

Setup of the adversary

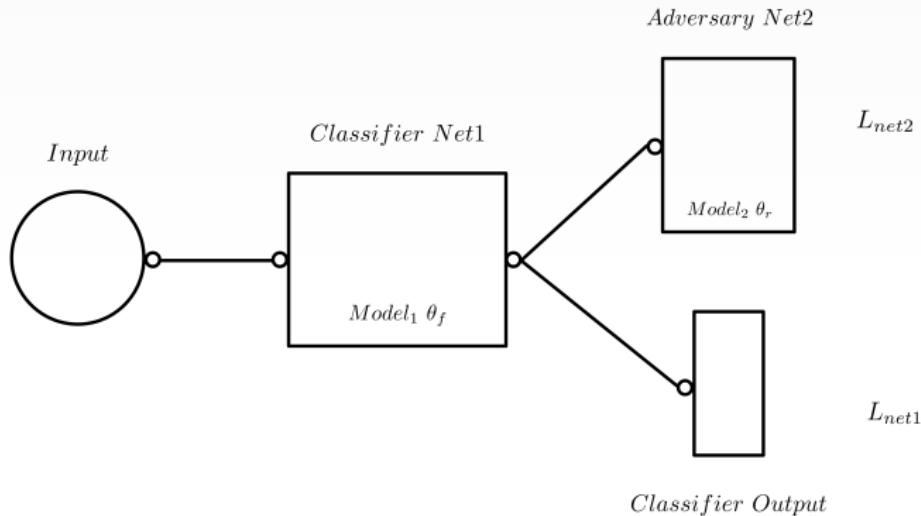
- Input: 14 variables motivated by a BDT variable scan.
- Hidden layers: 3 elu layers \times 32 nodes each
- Output layer: 1 sigmoid node
- Optimisation: SGD, Learning rate = 0.001, momentum = 0., no nesterov, no decay
- Duration: 400 iteration
- $\lambda = 0.1$

Approach 1 - Run 2 - Results



- Visible separation
- Adversary loss stops falling
- No improvement in the sensitivity visible

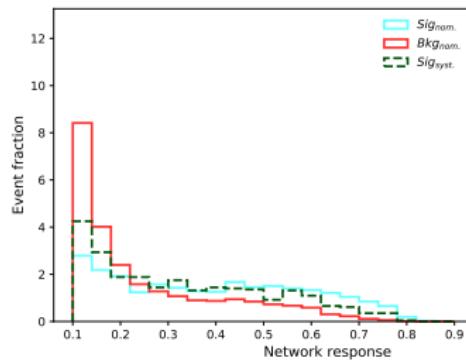
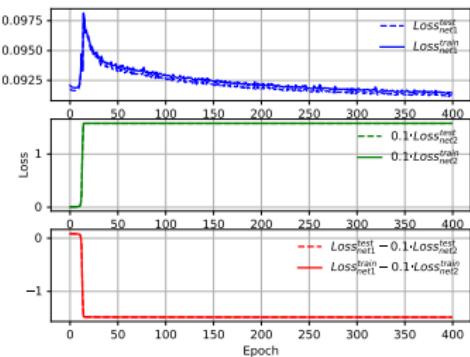
Approach 2



Input the last hidden layer to the adversary

- More information provided
- Possible danger: Missing last decision step

Approach 2 Results



- Losses behave as desired
- steep rise, although partially due to scale
- λ almost unusable due to difference in magnitude
- still high sensitivity

Approach 3

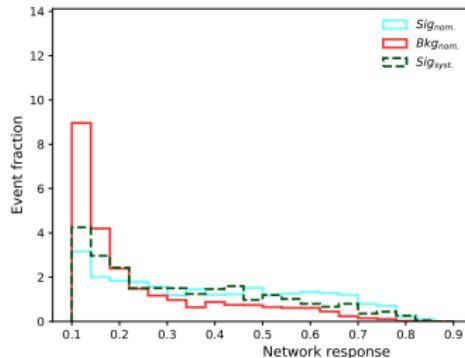
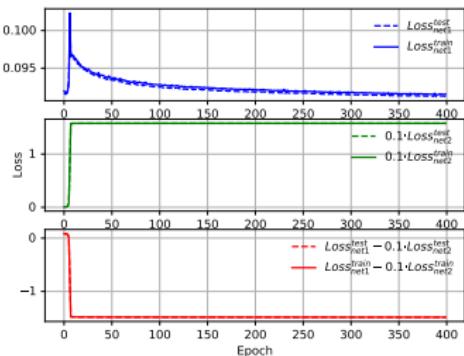
Assumption

Dimension of the hidden layer is too high

Easy workaround

- Add an intermediate layer
- Reduced to 16 nodes

Approach 3 Results



- Almost no difference
- Slightly higher stability

Thoughts

Input

- Different variable set → different problem topology
- Concatenate layers as input → Combination of more information and final output
- Separate background from the adversary's input → remove possible bias

General

- Different analysis → Insight on the network's general inability

Conclusions

Optimisation of a combination of networks

- An adversary cannot just be added
- An ANN has to be set up as a combination

Input of the adversary

- Input of the adversary does not justify the architecture
- Space for improvement in the input

Strong dependency on the learning rate

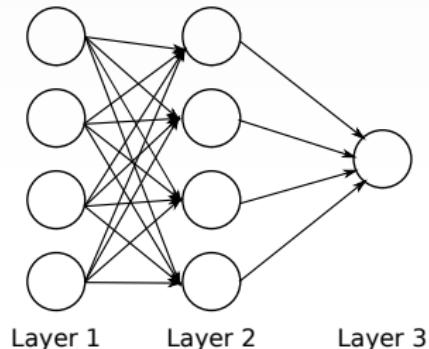
- Learning rate is the defining factor of the performance of the losses
- Presumably due to the fine topology around the minimum



<https://xkcd.com/1838/>

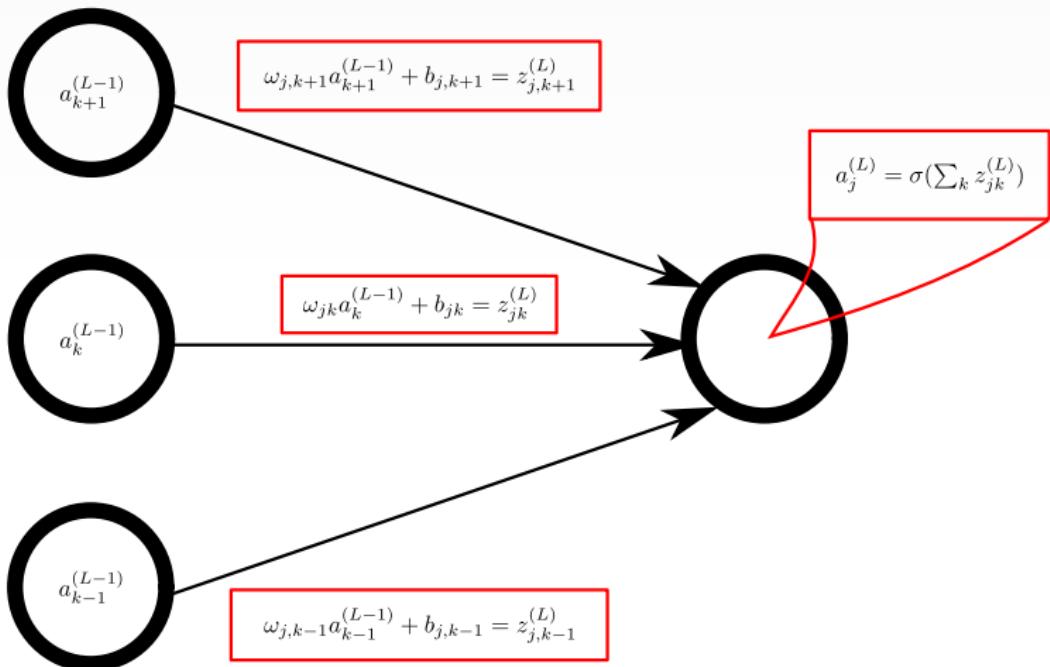
Neural Networks

- Built of layers of simple processors called nodes
- Each node is connected to each node in the surrounding layers
- The connection is built by a linear function with a weight ω and a bias b and a non linear activation function σ
- Learning is accomplished by backpropagation and a step optimizer



$$z_j^L = \omega_{jk}^L a_k^{L-1} + b_k$$
$$a_j^L = \sigma(z_j^L)$$

Neural Networks



Forward Propagation in a neural network

Choosing the next step

- In supervised learning truth tagged training data allows to calculate a cost function to estimate a model's quality. For example crossentropy:

$$C = -(y \log p + (1 - y) \log(1 - p)) \quad (1)$$

- Backpropagation estimates the parameters' impact on the cost function using the partial derivatives

$$\frac{\partial C}{\partial a_k^{L-1}} = \sum_{j=1}^N \frac{\partial z_j^L}{\partial a_k^{L-1}} \frac{\partial a_j^L}{\partial z_j^L} \frac{\partial C}{\partial a_j^L}$$

- Each parameter is then updated according to its impact on the cost function

Adversarial Neural Networks

- Originally introduced as Generative adversarial neural networks to overcome weaknesses of generative networks [2]
- Adds a second network controlling the dependency on features with large systematic uncertainties.
- The adversarial structure of the networks creates a minimax game
- Combined loss function

Optimisers

- Gradient Descent: Update parameters in the opposite direction of the error gradient
- Stochastic Gradient Descent: Gradient Descent but for each training example separately
- Adaptive Moment Estimation, Adam: adaptive learning rate based on exponentially decaying mean of past gradients and past square gradients.
- Others
- <https://towardsdatascience.com/types-of-optimization-algorithms-used-in-neural-networks-and-ways-to-optimize-gradient-95ae5d39529f>

Learning rate, momentum, decay

- Learning rate: Step-length in gradient direction

$$\theta' = \theta - lr \cdot g \quad (2)$$

- Momentum: create an adaptive and weight dependent learning rate

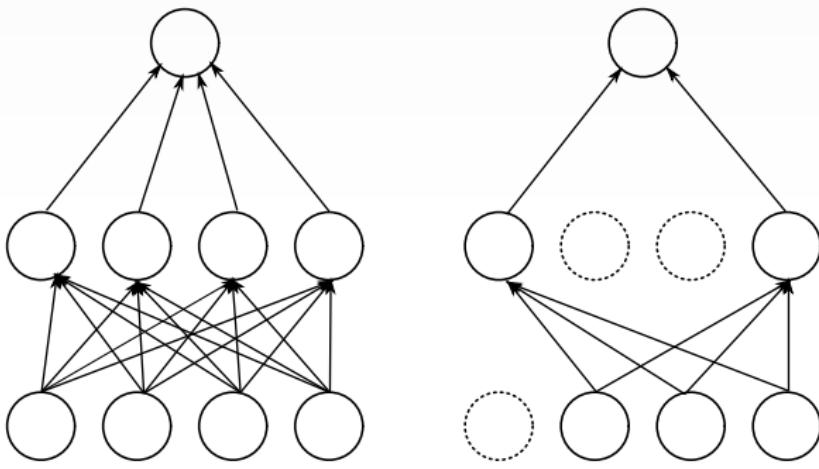
$$\nu' = \alpha\nu - \epsilon \frac{1}{m} \nabla_{\theta} \sum_j L(f(\hat{y}^j; \theta), y^j) \quad (3)$$

$$\theta' = \theta + \nu \quad (4)$$

- Decay: Decrease the learning rate over the course of the training

$$\epsilon' = \frac{\epsilon}{1 + \phi t} \quad (5)$$

Dropout



Without Dropout

With Dropout

Caption

Technical details

- The neural networks were built using Keras [1] with tensorflow [4] as backend